



Task	Classification	Reason
Study for exams	Collaborate with AI	AI explains concepts but I must learn them myself.
Build projects	Collaborate with AI	AI helps with coding but I make the decisions.
Debug code	Delegate to AI with review	AI finds bugs quickly; I verify the fix.
Write emails	Delegate to AI with review	AI improves grammar; I review before sending.
Update GitHub README	Delegate to AI with review	AI drafts documentation; I check accuracy.
Prepare for viva	Collaborate with AI	AI asks questions; I practice answering.
Learn new AI topics	Collaborate with AI	AI summarizes; I verify with official docs.
Plan weekly schedule	Fully automate	AI can generate reminders and schedules.
Decide career goals	Just me	This is a personal decision.
Write exams	Just me	AI should never take exams for me.



Hi Om Raj — good to meet you. I'll keep that learning style in mind for everything we work on together: step-by-step explanations, reasoning through concepts rather than just handing you finished answers, and clean, well-explained code when we get to Python/ML/MERN/Flutter/SQL work.

I've looked through the five PDFs you uploaded. They're all Anthropic's **AI Fluency course materials**, covering:

1. **The AI Fluency Framework (4D Model)** — four skills for working well with AI:
 - **Delegation** — deciding *whether, when, and how* to bring AI into a task
 - **Description** — describing your goal clearly so the AI gives useful output
 - **Discernment** — critically judging whether the AI's output is actually good
 - **Diligence** — taking responsibility for what you produce with AI's help
2. **Diligence, in more depth** — split into three parts: *Creation* (choosing tools thoughtfully), *Transparency* (disclosing AI's role), and *Deployment* (verifying/vouching for outputs before sharing them).
3. **Understanding Generative AI** — a primer on what makes generative AI different from traditional AI, the three enabling pillars (transformer algorithms, data, compute), how models are built (pre-training → fine-tuning → deployment), and current capabilities/limitations.

These read like slides from a course or workshop rather than a coding task — so before I dive in, let me check what you actually need from them.

What would you like help with on these AI Fluency documents? Explain the concepts in depth (step by step)


Let's build this in two layers: first the **technology** (what generative AI actually is), then the **skill** (how to work with it well). That mirrors how the documents themselves are structured, and it's the right learning order — you can't reason well about "should I delegate this to AI" until you understand what the AI is actually doing under the hood.

Part 1 — What generative AI actually is

The core distinction the slides draw is: traditional AI *classifies* or *scores* existing data (spam/not-spam), while generative AI *produces new* content it wasn't explicitly given. As a CS student, think of it as the difference between a discriminative function $f(x) \rightarrow \text{label}$ and a generative function $f(x) \rightarrow \text{new_sequence}$.

Three things had to come together for this to become practical:

- **Algorithms** — the 2017 transformer architecture. Its key trick is *self-attention*: instead of processing a sentence strictly left-to-right (like older RNNs), every token computes a relevance score against every other token in the sequence, in parallel. That's what let models handle long-range context efficiently and train on GPUs at scale.
- **Data explosion** — enough digitized text (web pages, code, books) to give the model billions of examples of "what word plausibly comes next."
- **Compute power** — GPUs made the matrix multiplications in attention layers cheap enough to run at massive scale.



Here's how a model actually gets built, stage by stage: If you've done any ML coursework, **pre-training** → **fine-tuning** should feel familiar — it's the same idea as transfer learning. Pre-training is like training a general-purpose feature extractor on a huge, generic dataset (self-supervised: the model just predicts the next token, over and over, across a huge corpus). Fine-tuning is the specialization step — like taking a pretrained CNN and retraining the final layers on your specific task, except here the "task" is *following instructions helpfully and safely* rather than one narrow label.

Capabilities that fall out of this: versatile language handling, general-purpose reasoning across domains, learning from examples given in the prompt (few-shot), and the ability to call external tools/APIs.

Limitations worth internalizing as an engineer:

- **Knowledge cutoff** — the model only "knows" what existed in its training data up to a certain date.
- **Hallucination** — it can generate fluent, confident, *wrong* text, because it's optimizing for "plausible next token," not "verified fact."
- **Context window** — there's a hard limit on how much text it can consider at once.
- **Weak spots in complex reasoning/math** — because it's fundamentally predicting tokens, not running a symbolic solver.


That last point is exactly *why* the AI Fluency framework exists — if the tool can be fluent and wrong at the same time, you need a discipline for using it well.

Part 2 — The AI Fluency framework

Before the 4 competencies, there's a foundational idea: **three modes of AI interaction**, which describe how much control you're handing over.

- **Automation** — you give a specific instruction, AI executes it. Like calling a function:
`format_code(x)`.
- **Augmentation** — you and the AI go back and forth, thinking together. Like pair programming.
- **Agency** — you configure the AI once, and it acts independently afterward, over multiple steps. Like setting up an autonomous coding agent to work through a backlog of tickets without you reviewing every step.

These aren't good/bad tiers — they're a dial you consciously choose, based on the stakes of the task. That choice *is* the first competency, **Delegation**.



Here's how the four competencies chain together in practice: Let's make that concrete with something you'd actually do — asking AI to help you write a sorting algorithm for an assignment:

Step	What it looks like
Delegation	You decide: is this a task I should hand to AI at all? Maybe not if the assignment is testing <i>your</i> understanding of recursion — but fine if you're using it to check your own implementation.
Description	You write a clear, specific prompt: "Here's my merge sort implementation, find any bugs in the merge step" — not just "fix my code."
Discernment	You don't blindly trust the output. You check: does the fix actually work? Did it change the time complexity? Does it match what I was taught?
Diligence	You take responsibility for the final code — you understand it well enough to explain it, and you're honest that AI helped if the assignment requires disclosure.

Notice **Discernment** is where most people get lazy — it's the step that directly fights hallucination. An AI can generate a sorting function that *looks* correct, compiles, and is still subtly wrong on an edge case. Discernment is basically "code review applied to AI output" — you'd never merge a teammate's PR without checking it; same standard applies here.

Diligence, in more depth

Diligence gets its own document because it's the "responsibility" layer, and it splits into three parts:

- **Creation diligence** — being deliberate about *which* AI tool you use and how you engage it. (Not every task needs the most powerful model, and not every task should use AI at all.)
- **Transparency diligence** — being honest about AI's role with whoever needs to know. For you as a student, this usually means: does your professor's policy require you to disclose AI use in an assignment? Different contexts — personal, academic, professional — have different disclosure expectations, and it's on you to know which one applies.
- **Deployment diligence** — before you submit, ship, or share AI-assisted work, *you* are vouching for it as if you wrote it yourself. If it's wrong, "the AI said so" isn't a defense — that's the whole point of this competency.

This is genuinely one of the more important ideas in here for you specifically: as AI coding assistants become normal in software engineering, **Discernment and Diligence are what separate a developer who uses AI well from one who just ships whatever it outputs.**

Want to test this by walking through Discernment/Diligence on a real piece of AI-generated code, or should I quiz you on the framework instead?